Experiment 3 · registered design, not results

Fixed comparison

50% credit / 50% generic

Same tables within each backbone

1,000 updates; 3 paired seeds

12 full-CPT baselines reused

48 additional runs

Settings fixed before benchmarking

Four additional methods

Full CPT without weight decay

L2-SP: anchor to released weights

DiagScale: query scales only

LoRA: rank-16 Q/V updates

Compare with ordinary full CPT

Report credit + OOD at 4 steps